Zhuofu Chen

aetiurf@gmail.com · chenzhuofu · https://chenzhuofu.github.io/ ·

Education

Tongji University Shanghai, China

B.S. in Computer Science and Technology (Elite Class) Sept. 2021 - Present

- GPA: 4.87/5.00 Ranking: 3/20

Research Interests

I have broad interests in building system infrastructures to systematically bring better *performance*, *scalability* and *usability* to *real-world applications*. Specifically, I often contemplate how to *redesign next-generation datacenter systems* to bridge the gap between evolving hardware and emerging needs of software, and to serve data-intensive applications such as *AI inference and training*.

Experience

Catalyst Group, Carnegie Mellon University

Research Intern advised by Zhihao Jia May. 2024 - Present

- Optimized the LLM serving system for large batch decoding, achieving 7x latency improvement.
- Designed a request level scheduler considering different SLO requirements across requests, achieving up to 73% higher SLO attainment.
- Implemented a novel position-persistent mechanism to reduce the overhead of sparse attention, demonstrating 20% improvement over the state-of-the-art.
- Two papers are pending review.

Institute of Parallel and Distributed Systems, Shanghai Jiao Tong University

Research Intern advised by Xingda Wei and Rong Chen Jul. 2023 - Apr. 2024

- Developed an optimal GPU disaggregation system for transparently and efficiently serving AI applications with no performance degradation compared to local execution.
- Created a theoretical model characterizing the overhead of disaggregating applications, achieving an error margin within 10%.
- Built a profiling-based tool that allows users to derive network requirements for any AI applications.
- One paper is pending review.

Key Laboratory of Embedded System and Service Computing, Tongji University

Research Intern advised by Zhijun Ding Nov. 2022 - Oct. 2023

- Proposed a dynamic import mechanism to enhance WebAssembly-based runtime, introducing a shareable resource allocation paradigm.
- Implemented a state transfer method through shared memory, improving serialization into a storage layer.
- Implemented an OCI shim to bridge orchestration tools and runtime.

Publications

(* indicates equal contribution)

[1] Tianxia Wang*, **Zhuofu Chen***, Xingda Wei, Jinyu Gu, Rong Chen, Haibo Chen. Characterizing Network Requirements for GPU API Remoting in AI Applications, *Under review.*

[2] Zikun Li*, **Zhuofu Chen***, Remi Delacourt, Gabriele Oliaro, Zeyu Wang, Qinghan Chen, Shuhuai Lin, April Yang, Zhihao Zhang, Zhuoming Chen, Sean Lai, Xupeng Miao, Zhihao Jia. AdaServe: SLO-Customized LLM Serving with Fine-Grained Speculative Decoding, *Under review.*

[3] Lijie Yang*, Zhihao Zhang*, **Zhuofu Chen**, Zikun Li, Zhihao Jia. TidalDecode: Fast and Accurate LLM Decoding with Position Persistent Sparse Attention, *International Conference on Learning Representations (ICLR), 2025.*

Selected Projects

XPURemoting

- A GPU disaggregation system redirecting CUDA API calls to remote devices with near-native performance.
- An easy-to-use perf tool to model the overhead of disaggregation for arbitrary applications.
- 15k lines of code in Rust.

FlexFlow SpecScheduler

- A performant LLM serving system for large batch decoding, especially for speculative decoding.
- An request-level scheduler aiming to maximize SLO attainment across diverse tasks.
- 10k lines of code in C++/CUDA.

Selected Awards

SOSP Student Scholarship	2024
Tongji University Lu Hao Scholarship	2024
National 1st Prize (0.55%) in Contemporary Undergraduate Mathematical Contest in Modeling	2023
Regional 1st Prize in Contemporary Undergraduate Mathematical Contest in Modeling	2023
China National Scholarship (top 0.2%)	2022
Regional 2nd Prize in Contemporary Undergraduate Mathematical Contest in Modeling	2022
1st Prize of National Olympiad in Informatics in Provinces	2019